

Complete Project Guide: Understanding Every Component

Project Overview

Your project is a **modular RL research framework** designed to test different combinations of algorithms (PPO, Domain Randomization, SR2L) on a quadruped robot. Think of it as a scientific experiment where you're testing 4 different "recipes" to see which makes the most robust robot.

Your 4 Experiments (Ablations):

1. PPO only (baseline)
2. PPO + Domain Randomization (DR)
3. PPO + Smooth Regularization (SR2L)
4. PPO + DR + SR2L (full method)

Directory Structure Explained

1. configs/ - The Brain of Your Experiments

```
configs/
├── experiments/    # Your 4 main experiments
│   ├── ppo_baseline.yaml    # Experiment 1: Just PPO
│   ├── ppo_dr.yaml         # Experiment 2: PPO + DR
│   ├── ppo_sr2l.yaml       # Experiment 3: PPO + SR2L
│   └── ppo_dr_sr2l.yaml    # Experiment 4: Everything
├── env/            # Environment settings
│   └── realant.yaml     # Robot configuration
├── train/          # Training settings
│   └── default.yaml     # Base hyperparameters
├── eval/           # Evaluation settings
│   └── default.yaml     # How to test your robot
```

What you'll modify:

- Hyperparameters in `train/default.yaml`
- Fault settings in experiment files
- Environment rewards in `env/realant.yaml`

2. src/ - The Code That Runs Everything

```

src/
├── agents/          # RL algorithms
│   ├── ppo.py       # Standard PPO implementation
│   └── sr2l_ppo.py   # PPO + SR2L (smooth regularization)
├── envs/            # Robot environment
│   ├── realant_env.py # Main robot environment
│   ├── fault_injection.py # Breaks joints, adds noise
│   └── curriculum.py  # Gradually increases difficulty
├── utils/           # Helper tools
│   ├── config.py     # Loads YAML configs
│   ├── logger.py     # Logs to W&B, TensorBoard
│   └── metrics.py     # Calculates success rates
└── train.py         # Main script you'll run

```

What you'll modify:

- Add SR2L logic to `sr2l_ppo.py`
- Implement fault scenarios in `fault_injection.py`
- Define curriculum phases in `curriculum.py`

3. experiments/ - Where Results Live

```

experiments/
├── ppo_dr_sr2l_2025_07_30_143052/ # Auto-generated folder
│   ├── logs/                      # Training logs
│   │   ├── progress.csv           # Raw numbers
│   │   └── tensorboard/           # TB files
│   ├── models/                   # Saved neural networks
│   │   ├── checkpoint_1000.pt     # Early model
│   │   └── best_model.pt          # Best performance
│   ├── videos/                   # Robot recordings
│   └── metrics/                  # Evaluation results
│       └── evaluation_results.json # Final scores

```

What happens here:

- Automatically created for each run
- All results saved with timestamps
- Easy to compare experiments

Phase 1: Local Development (Your MacBook)

mermaid

```
graph LR
  A[Write Code] --> B[Test Locally]
  B --> C{Works?}
  C -->|No| A
  C -->|Yes| D[Sync to Cluster]
```

1. **Edit code** in `src/`
2. **Run quick test:**

bash

```
python src/train.py --config configs/experiments/ppo_baseline.yaml \
  --override train.total_timesteps=1000 \
  --override train.num_envs=1
```

3. **Check it doesn't crash** (takes ~1 minute)

Phase 2: Cluster Training (Full Experiments)

mermaid

```
graph LR
  A[Sync Code] --> B[Submit Job]
  B --> C[Train 10M Steps]
  C --> D[Monitor W&B]
  D --> E[Download Results]
```

1. **Sync to cluster:**

bash

```
./scripts/sync_to_cluster.sh
```

2. **Submit SLURM job:**

bash

```
sbatch scripts/train_cluster.sh ppo_dr_sr2l
```

3. Monitor on Weights & Biases (real-time)

Weights & Biases Setup

What is W&B?

- Cloud-based experiment tracking
- Real-time training curves
- Hyperparameter comparison
- Video recordings of your robot

Setup Steps:

1. **Create account** at wandb.ai
2. **Install and login:**

```
bash

pip install wandb
wandb login # Enter your API key
```

3. **Enable in config:**

```
yaml

# In configs/train/default.yaml
logging:
  wandb: true # Change from false
  wandb_project: "robust-quadruped"
  wandb_entity: "your-username"
```

What W&B Tracks:

- Learning curves (reward over time)
- Success rates
- Videos of robot walking
- All hyperparameters
- System metrics (GPU usage)

Your Experimental Pipeline

Step 1: Baseline Test (PPO Only)

```
bash

# Local quick test
python src/train.py --config configs/experiments/ppo_baseline.yaml \
    --override train.total_timesteps=10000

# Cluster full run
sbatch scripts/train_cluster.sh ppo_baseline
```

Step 2: Add Domain Randomization

```
bash

# This config enables joint failures and sensor noise
sbatch scripts/train_cluster.sh ppo_dr
```

Step 3: Add Smooth Regularization

```
bash

# This config enables SR2L for smooth actions
sbatch scripts/train_cluster.sh ppo_sr2l
```

Step 4: Full Method

```
bash

# Everything combined
sbatch scripts/train_cluster.sh ppo_dr_sr2l
```

Step 5: Evaluation

```
bash

# Test all trained models
python scripts/evaluate_all.py --fault_scenarios all
```

Key Files You'll Work With

1. src/train.py - Main Training Loop

python

Pseudocode of what happens

def main():

1. Load config

config = load_config(args.config)

2. Create environment

env = create_env(config)

if config.domain_randomization.enabled:

env = add_fault_wrapper(env)

3. Create agent

if config.sr2l.enabled:

agent = SR2L_PPO(env, config)

else:

agent = PPO(env, config)

4. Train

for step **in** range(config.total_timesteps):

agent.train()

if step % config.eval_freq == 0:

evaluate_and_log()

2. src/envs/fault_injection.py - Breaking the Robot

python

```

class FaultInjectionWrapper:
    def __init__(self, env, config):
        self.fault_config = config.faults

    def reset(self):
        # Randomly select which joints to break
        self.failed_joints = random.sample(
            range(8),
            k=random.randint(0, self.max_failures)
        )

    def step(self, action):
        # Override actions for broken joints
        action[self.failed_joints] = 0

        # Add sensor noise
        obs += np.random.normal(0, self.sensor_noise_std)

```

3. src/agents/sr2l_ppo.py - Smooth Regularization

python

```

class SR2L_PPO(PPO):
    def compute_loss(self, obs, actions):
        # Standard PPO loss
        ppo_loss = super().compute_loss(obs, actions)

        # SR2L: Penalize non-smooth policies
        perturbed_obs = obs + small_noise
        smooth_loss = MSE(
            self.policy(obs),
            self.policy(perturbed_obs)
        )

        return ppo_loss + self.sr2l_lambda * smooth_loss

```



Understanding Your Results

Training Curves (W&B)

- **Episode Reward:** Should increase over time
- **Success Rate:** % of episodes robot walks successfully

- **Policy Loss:** Should decrease and stabilize
- **Value Loss:** Indicates learning quality

Evaluation Metrics

```
json
{
  "clean_performance": {
    "success_rate": 0.95,
    "avg_distance": 4.8,
    "avg_reward": 245.6
  },
  "single_joint_failure": {
    "success_rate": 0.78,
    "recovery_time": 1.2
  },
  "multiple_failures": {
    "success_rate": 0.45,
    "failure_modes": ["fall", "spin", "stuck"]
  }
}
```

Typical Development Cycle

Week 1-2: Setup & Baseline

1. Get PPO baseline working
2. Verify environment runs correctly
3. Establish baseline metrics

Week 3-4: Implement DR

1. Add fault injection wrapper
2. Implement curriculum learning
3. Test with increasing fault rates

Week 5-6: Implement SR2L

1. Modify PPO loss function
2. Tune smoothness parameter
3. Verify it doesn't hurt performance

Week 7-8: Full Experiments

1. Run all 4 ablations on cluster
2. Each takes ~24-48 hours
3. Monitor via W&B

Week 9-10: Analysis

1. Statistical comparison
2. Generate plots
3. Failure mode analysis

Quick Reference Commands

bash

Local testing (1 min)

```
python src/train.py --config configs/experiments/ppo_baseline.yaml \  
    --override train.total_timesteps=1000
```

View tensorboard locally

```
tensorboard --logdir experiments/
```

Sync to cluster

```
rsync -av --exclude='experiments/' --exclude='venv/' \  
    ./ cluster:~/projects/robust-quadruped/
```

Submit cluster job

```
ssh cluster
```

```
cd projects/robust-quadruped
```

```
sbatch scripts/train_cluster.sh ppo_dr_sr2l
```

Monitor W&B

Just go to wandb.ai/your-username/robust-quadruped

Download results

```
rsync -av cluster:~/projects/robust-quadruped/experiments/ ./experiments/
```

Run evaluation

```
python scripts/evaluate_all.py --models experiments/*/models/best_model.pt
```

? Common Issues & Solutions

"Environment not found"

- Make sure MuJoCo is installed
- Check RealAnt package is installed

"CUDA out of memory"

- Reduce `num_envs` in config
- Reduce `batch_size`

"Training unstable"

- Lower learning rate
- Increase `n_epochs`
- Check reward scale

"W&B not logging"

- Check you're logged in: `wandb login`
- Verify `wandb: true` in config
- Check internet connection



What to Learn

1. Hydra Configuration

- [Tutorial](#)
- Handles YAML loading and overrides

2. Stable-Baselines3

- [Docs](#)
- Your PPO implementation base

3. Weights & Biases

- [Quickstart](#)
- Focus on: logging metrics, saving videos

4. MuJoCo Basics

- Understanding observation/action spaces
- Coordinate systems



Success Criteria

Your project succeeds when:

1. **All 4 ablations train successfully**
2. **Clear performance ordering:** Full > DR/SR2L > Baseline
3. **Robust to failures:** >70% success with 1 joint failure
4. **Smooth behaviors:** No jerky movements with SR2L
5. **Reproducible:** Same config = similar results

Pro Tips

1. **Start simple:** Get baseline working first
2. **Version control:** Commit before each experiment
3. **Document bugs:** Keep a log of issues/solutions
4. **Visualize early:** Watch videos of robot behavior
5. **Incremental testing:** Test each component separately
6. **Use W&B:** It's much easier than managing files manually
7. **Ask for help:** When stuck for >2 hours

This is your complete roadmap! Start with the baseline, gradually add complexity, and always test locally before cluster runs.