

Step-by-Step Command Guide

Day 1: Initial Setup

1. Create the project structure

```
bash

# Download and run the setup script
python setup_project.py --path robust-quadruped-rl
cd robust-quadruped-rl
```

2. Set up Python environment

```
bash

# Create virtual environment
python3 -m venv venv

# Activate it (Mac/Linux)
source venv/bin/activate

# Install dependencies
pip install --upgrade pip
pip install -r requirements.txt
pip install -e .
```

3. Set up Weights & Biases

```
bash

# Install and login
pip install wandb
wandb login
# Enter your API key from https://wandb.ai/authorize
```

4. Test installation

```
bash

# Test imports work
python -c "import stable_baselines3; import mujoco; import wandb; print('All good!')"
```

Day 2-3: Implement Basic Components

1. Create the main training script

```
bash

# Create src/train.py
# This is where you'll implement the training loop
```

2. Test with minimal config

```
bash

# Run for just 1000 steps to check it works
python src/train.py \
  --config configs/experiments/ppo_baseline.yaml \
  --override train.total_timesteps=1000 \
  --override train.num_envs=1 \
  --override logging.wandb=false
```

3. Check output structure

```
bash

# Should create:
ls experiments/
# ppo_baseline_2025_07_30_143052/

# Check logs exist:
ls experiments/ppo_baseline_*/logs/
# progress.csv
```

Day 4-5: Local Testing Phase

1. Test baseline PPO (no faults)

```
bash
```

```
# 10K steps, ~5 minutes on Mac
python src/train.py \
  --config configs/experiments/ppo_baseline.yaml \
  --override train.total_timesteps=10000 \
  --override train.num_envs=2 \
  --override logging.wandb=true \
  --override logging.wandb_project="robust-quadruped-test"
```

2. View results locally

```
bash

# Start TensorBoard
tensorboard --logdir experiments/

# Open browser to http://localhost:6006
```

3. Test with faults enabled

```
bash

# Quick test of domain randomization
python src/train.py \
  --config configs/experiments/ppo_dr.yaml \
  --override train.total_timesteps=5000 \
  --override train.num_envs=1
```

Day 6: Cluster Setup

1. Sync code to cluster

```
bash
```

```
# Create sync script
cat > scripts/sync_to_cluster.sh << 'EOF'
#!/bin/bash
CLUSTER_USER="your_username"
CLUSTER_HOST="cluster.university.edu"
REMOTE_DIR="/projects/robust-quadruped"

rsync -avz --exclude='venv/' --exclude='experiments/' --exclude='__pycache__/' \
  ./ ${CLUSTER_USER}@${CLUSTER_HOST}:${REMOTE_DIR}/
EOF

chmod +x scripts/sync_to_cluster.sh
./scripts/sync_to_cluster.sh
```

2. Create SLURM job script

```
bash

cat > scripts/train_cluster.sh << 'EOF'
#!/bin/bash
#SBATCH --job-name=robust_rl
#SBATCH --output=logs/%x_%j.out
#SBATCH --error=logs/%x_%j.err
#SBATCH --time=48:00:00
#SBATCH --nodes=1
#SBATCH --gpus-per-node=1
#SBATCH --cpus-per-task=8
#SBATCH --mem=32G

# Load modules
module load python/3.9
module load cuda/11.8

# Activate environment
source venv/bin/activate

# Run training
python src/train.py --config configs/experiments/$1.yaml

EOF
```

3. SSH to cluster and setup

```
bash
```

```
ssh cluster
```

```
cd ~/projects/robust-quadruped
```

```
module load python/3.9
```

```
python -m venv venv
```

```
source venv/bin/activate
```

```
pip install -r requirements.txt
```

Day 7-14: Full Training Runs

1. Submit baseline job

```
bash
```

```
# On cluster
```

```
sbatch scripts/train_cluster.sh ppo_baseline
```

```
# Note the job ID (e.g., 12345)
```

2. Monitor job

```
bash
```

```
# Check job status
```

```
squeue -u $USER
```

```
# Watch logs
```

```
tail -f logs/robust_rl_12345.out
```

```
# Check W&B dashboard
```

```
# Go to: https://wandb.ai/your-username/robust-quadruped
```

3. Submit all ablations

```
bash
```

```
# Run all 4 experiments
```

```
for exp in ppo_baseline ppo_dr ppo_sr2l ppo_dr_sr2l; do
```

```
    sbatch scripts/train_cluster.sh $exp
```

```
    sleep 10 # Avoid overwhelming scheduler
```

```
done
```

Day 15-16: Evaluation

1. Download trained models

```
bash

# From your local machine
rsync -avz cluster:~/projects/robust-quadruped/experiments/ ./experiments/
```

2. Run evaluation script

```
bash

# Create evaluation script first
python scripts/evaluate_all.py \
    --models experiments/*/models/best_model.pt \
    --output results/evaluation_results.csv
```

3. Generate plots

```
bash

# Open Jupyter notebook
jupyter notebook notebooks/visualize_results.ipynb
```

Common Tasks

Check GPU usage on cluster

```
bash

nvidia-smi
```

Kill a running job

```
bash

scancel <job_id>
```

Debug a failed run

```
bash
```

```
# Check error logs
cat logs/robust_rl_12345.err
```

```
# Common fixes:
# - Out of memory: reduce batch_size or num_envs
# - Module not found: check virtual env is activated
# - CUDA error: check GPU is allocated
```

Resume from checkpoint

```
bash

python src/train.py \
  --config configs/experiments/ppo_dr.yaml \
  --resume experiments/ppo_dr_2025_07_30_143052/models/checkpoint_5000000.pt
```

Quick visualization

```
bash

# Record video of trained policy
python scripts/record_video.py \
  --model experiments/ppo_dr_sr2l_*/models/best_model.pt \
  --num_episodes 5 \
  --fault_scenario "single_joint"
```



W&B Specific Commands

Create new project

```
python
```

```
# In your train.py
import wandb

wandb.init(
    project="robust-quadruped",
    name=f"{config.experiment.name}_{timestamp}",
    config=config,
    tags=["ablation", config.experiment.name]
)

# Log metrics
wandb.log({
    "episode_reward": episode_reward,
    "success_rate": success_rate,
    "learning_rate": current_lr,
})

# Log video
wandb.log({"video": wandb.Video(video_path, fps=30)})
```

Compare runs

```
bash

# Use W&B dashboard to:
# 1. Select multiple runs
# 2. Create comparison plots
# 3. Download data as CSV
```

Quick Testing Cheatsheet

```
bash
```


Super quick test (1 min)

```
python src/train.py --config configs/experiments/ppo_baseline.yaml \  
  --override train.total_timesteps=1000 \  
  --override train.num_envs=1 \  
  --override logging.wandb=false
```

Medium test (10 min)

```
python src/train.py --config configs/experiments/ppo_dr.yaml \  
  --override train.total_timesteps=50000 \  
  --override train.num_envs=4
```

Full local test (1 hour)

```
python src/train.py --config configs/experiments/ppo_dr_sr2l.yaml \  
  --override train.total_timesteps=500000
```

Cluster production run (24-48 hours)

```
sbatch scripts/train_cluster.sh ppo_dr_sr2l
```



Progress Checklist

- ☐ Project structure created
- ☐ Virtual environment set up
- ☐ W&B account created and logged in
- ☐ Basic PPO training works locally
- ☐ TensorBoard shows training curves
- ☐ Cluster access configured
- ☐ First cluster job submitted successfully
- ☐ All 4 ablations submitted
- ☐ Models downloaded from cluster
- ☐ Evaluation results generated
- ☐ Plots and analysis complete

Remember: Start simple, test often, and gradually increase complexity!